

VISVESVARAYA TECHNOLOGICAL UNIVERSITY
“JnanaSangama”, Belgaum -590014, Karnataka.



LAB REPORT
ON
MACHINE LEARNING

Submitted by:

HARSHITHA R (1BM21CS075)

in partial fulfillment for the award of the degree of

BACHELOR OF ENGINEERING
in
COMPUTER SCIENCE AND ENGINEERING



B.M.S. COLLEGE OF ENGINEERING
(Autonomous Institution under VTU)
BENGALURU-560019
March 2024-June 2024



**B. M. S. College of Engineering,
Bull Temple Road, Bangalore 560019**
(Affiliated To Visvesvaraya Technological University, Belgaum)
Department of Computer Science and Engineering

CERTIFICATE

This is to certify that the Lab work entitled “**MACHINE LEARNING**” carried out by **HARSHITHA R(1BM21CS075)**, who is bonafide student of **B. M. S. College of Engineering**. It is in partial fulfillment for the award of **Bachelor of Engineering in Computer Science and Engineering** of the Visvesvaraya Technological University, Belgaum during the year 2023-24. The Lab report has been approved as it satisfies the academic requirements in respect of Machine Learning Lab - **(22CS3PCMAL)** work prescribed for the said degree.

Dr. K. Panimozhi

Assistant Professor
Department of CSE
BMSCE, Bengaluru

Dr. Jyothi S Nayak

Professor and Head
Department of CSE
BMSCE, Bengaluru

INDEX

Sl. No.	Experiment Title	Page No.
1	Write a python program to import and export data using Pandas library functions	3-5
2	Demonstrate various data pre-processing techniques for a given dataset	6-11
3	Use an appropriate data set for building the decision tree (ID3) and apply this knowledge to classify a new sample.	12-18
4	Implement Linear and Multi-Linear Regression algorithm using appropriate dataset	19-25
5	Build Logistic Regression Model for a given dataset	26-30
6	Build KNN Classification model for a given dataset.	31-34
7	Build Support vector machine model for a given dataset	35-38
8	Build Artificial Neural Network model with back propagation on a given dataset	39-41
9	a)Implement Random forest ensemble method on a given dataset. b)Implement Boosting ensemble method on a given dataset.	42-51
10	Build k-Means algorithm to cluster a set of data stored in a .CSV file.	52-55
11	Implement Dimensionality reduction using Principle Component Analysis (PCA) method.	56-59

PROGRAM 1

Date:05-04-2024

Write a python program to import and export data using Pandas library functions

05/04/24

Program -1

Write a python program to import and export data using Pandas library functions

IMPORT:

```
import pandas as pd  
airbnb_data = pd.read_csv("listings.csv")  
airbnb_data.head()
```

EXPORT:

```
airbnb_data.to_csv("exported-listings.csv")
```

READING DATA FROM URL:

```
url = "https://archive.ics.uci.edu/ml/machine-learning-  
databases/iris/iris.data"  
col_names = ["sepal-length-in-cm", "sepal-width-  
in-cm", "petal-length-in-cm",  
"petal-width-in-cm", "class"]  
iris_data = pd.read_csv(url, names=col_names)  
iris_data.head()
```


Import:

```
import pandas as pd  
# Read the CSV file  
airbnb_data = pd.read_csv("listings.csv")  
# View the first 5 rows  
airbnb_data.head()
```

[1]:		id	name	host_id	host_name	neighbourhood_group	neighbourhood	latitude	longitude	room_type	price	minimum_nights	number_of_reviews	last_review
0	329172	Hillside designer home,10 min.dwntn	1680871	Janet		Nan	78746	30.30085	-97.80794	Entire home/apt	495	3	7	2022-
1	329306	Urban Homestead, 5 minutes to downtown	880571	Angel		Nan	78702	30.27232	-97.72579	Private room	63	2	570	2022-
2	331549	One Room with Private Bathroom	1690383	Sandra		Nan	78725	30.23911	-97.58625	Private room	100	2	0	
3	333815	Solar Sanctuary - Austin Room	372962	Kim		Nan	78704	30.25381	-97.75262	Private room	102	2	164	2022-
4	333442	Rare Secluded 1940s Estate	1698318	Virginia		Nan	78703	30.31267	-97.76641	Entire home/apt	286	3	163	2022-

Export:

```
airbnb_data.to_csv("exported_listings.csv")
```



A	B	C	D	E	F	G	H	I	J	K	L	M	N	O	P	Q	R	S
1																		
2	0	229172 Hillside dr	1680871	Janet		78746	30.30085	-97.80794	Entire home	495	3	7	#####	0.05	1	363	1	
3	1	329306 Urban Hom	880571	Angel		78702	30.27232	-97.72579	Private ro	63	2	570	#####	4.36	5	55	45	
4	2	331549 One Room	1690383	Sandra		78725	30.23911	-97.58625	Private room	100	2	0	1	0	0	0	0	
5	3	333815 Solar Sanctu	372962	Kim		78704	30.25381	-97.75262	Private ro	102	2	164	#####	1.26	1	36	18	
6	4	333442 Rare Seclu	1698318	Virginia		78703	30.31267	-97.76641	Entire ho	286	3	163	#####	1.32	1	271	15	
7	5	335885 Abcd/2bar	1707903	Stephen		78749	30.20621	-97.8558	Entire ho	300	28	43	#####	0.33	1	90	0	
8	6	334616 Great SKS	608539 D			78741	30.24481	-97.7238	Entire ho	250	1	2	#####	0.02	1	0	0	
9	7	337125 1800 Sq ft	261883	Carolyn		78759	30.42035	-97.7699	Private ro	80	32	33	#####	0.25	1	83	0	
10	8	340045 Greenbelt	1725752	Shanti		78746	30.25937	-97.7925	Entire ho	400	1	1	#####	0.09	1	0	1	
11	9	335945 The Sheri	1555683	Patrick		78741	30.24026	-97.7334	Entire hor	221	3	34	#####	0.26	1	27	2	
12	10	340630 2 Bedroom	531267	Shelley		78757	30.33061	-97.7253	Entire hor	50	30	23	#####	0.18	3	25	3	
13	11	341596 South Aus	1733004	Julie		78704	30.24623	-97.759	Private ro	120	2	9	#####	0.69	1	122	15	
14	12	341382 SXSW 41	1205884	Luke And Rach		78702	30.26543	-97.7134	Entire hor	750	4	12	#####	0.95	1	0	0	
15	13	342243 Spacious L	1736662	Teresa		78759	30.41633	-97.7523	Private ro	45	31	3	#####	0.06	1	317	0	
16	14	342039 Garden Gu	1735494	Steph And Joey		78702	30.25158	-97.7315	Entire hor	116	1	342	#####	2.62	2	120	30	
17	15	343889 Close-In C	1744639	Darcy		78702	30.27293	-97.7229	Entire hor	47	30	32	#####	0.24	1	214	4	
18	16	343462 Charming	1742384	Rachel		78721	30.26895	-97.6895	Entire hor	99	1	255	#####	1.95	1	143	12	
19	17	345473 Exquisite	1751224	Whitney		78704	30.24619	-97.7457	Entire hor	308	2	164	#####	1.38	1	86	44	
20	18	345221 Austin His	1644218	Cecily		78701	30.27247	-97.748	Entire hor	299	3	8	#####	0.06	2	0	0	
21	19	3456 Walk to 6t	8028	Sylvia		78702	30.26057	-97.7344	Entire hor	96	2	623	#####	3.71	1	318	44	
22	20	347572 East Austl	1761506	Savanna		78702	30.26448	-97.7172	Entire hor	125	2	29	#####	0.22	1	0	0	
23	21	34769 NW Austin	8186	Elizabeth		78729	30.45697	-97.7842	Private ro	41	1	273	#####	1.77	1	14	9	
24	22	345536 Artist's Hn	1752493	Gigi		78704	30.2466	-97.7616	Entire hor	156	3	284	#####	2.19	2	319	29	
25	23	347935 for formul	1763742	Amber		78704	30.25152	-97.7621	Entire hor	643	5	1	#####	0.01	1	0	0	

Reading data from URL:

```
url = "https://archive.ics.uci.edu/ml/machine-learning-databases/iris/iris.data"

# Define the column names
col_names = ["sepal_length_in_cm",
             "sepal_width_in_cm",
             "petal_length_in_cm",
             "petal_width_in_cm",
             "class"]

# Read data from URL
iris_data = pd.read_csv(url, names=col_names)

iris_data.head()
```

[10]:

	sepal_length_in_cm	sepal_width_in_cm	petal_length_in_cm	petal_width_in_cm	class
0	5.1	3.5	1.4	0.2	Iris-setosa
1	4.9	3.0	1.4	0.2	Iris-setosa
2	4.7	3.2	1.3	0.2	Iris-setosa
3	4.6	3.1	1.5	0.2	Iris-setosa
4	5.0	3.6	1.4	0.2	Iris-setosa

PROGRAM 2

Date:05-04-2024

Demonstrate various data pre-processing techniques for a given dataset

Code and Output

2. Importing and Exploration of the dataset

```
In [2]: # Loading the data and setting the unique client_id as the index::  
df = pd.read_csv('/content/loans.csv', index_col = 'client_id')
```

```
In [3]: ## showing the first 5 rows of the dataset:  
df.head()
```

```
Out[3]:
```

	loan_type	loan_amount	repaid	loan_id	loan_start	loan_end	rate
client_id							
46109	home	13672	0	10243	2002-04-16	2003-12-20	2.15
46109	credit	9794	0	10984	2003-10-21	2005-07-17	1.25
46109	home	12734	1	10990	2006-02-01	2007-07-05	0.68
46109	cash	12518	1	10596	2010-12-08	2013-05-05	1.24
46109	credit	14049	1	11415	2010-07-07	2012-05-21	3.13

```
In [4]: # To check the Dimensions of the dataset:  
df.shape
```

```
Out[4]: (443, 7)
```

```
In [5]: # Checking the info of the data:  
df.info()
```

```
<class 'pandas.core.frame.DataFrame'>  
Index: 443 entries, 46109 to 26945  
Data columns (total 7 columns):  
 #   Column      Non-Null Count  Dtype     
---    
 0   loan_type    443 non-null   object    
 1   loan_amount  443 non-null   int64    
 2   repaid       443 non-null   int64    
 3   loan_id      443 non-null   int64    
 4   loan_start   443 non-null   object    
 5   loan_end     443 non-null   object    
 6   rate         443 non-null   float64  
dtypes: float64(1), int64(3), object(3)  
memory usage: 27.7+ KB
```

3. Checking the datatypes of the columns

In [6]:

```
df.dtypes
```

```
Out[6]: loan_type      object
loan_amount     int64
repaid         int64
loan_id        int64
loan_start     object
loan_end       object
rate           float64
dtype: object
```

4. Converting the data types of columns

- loan_id to object
- repaid to category dtype
- loan_start and loan_end to date type

In [7]:

```
# Loan_id:
df['loan_id'] = df['loan_id'].astype('object')

# repaid:
df['repaid'] = df['repaid'].astype('category')
```

In [8]:

```
# Loan_start:
df['loan_start'] = pd.to_datetime(df['loan_start'], format = '%Y-%m-%d')

# Loan_end:
df['loan_end'] = pd.to_datetime(df['loan_end'], format = '%Y-%m-%d')
```

Checking the datatypes again:

In [9]:

```
df.dtypes
```

```
Out[9]: loan_type      object
loan_amount     int64
repaid         category
loan_id        object
loan_start     datetime64[ns]
loan_end       datetime64[ns]
rate           float64
dtype: object
```

5. Summary Statistics of the data

In [10]:

```
# Summary Statistics for Numerical data:
df.describe()
```

Out[10]:

	loan_amount	loan_start	loan_end	rate
count	443.000000	443	443	443.000000
mean	7982.311512	2007-08-02 12:56:53.092550912	2009-08-23 11:35:37.246049536	3.217156
min	559.000000	2000-01-26 00:00:00	2001-08-02 00:00:00	0.010000
25%	4232.500000	2003-10-19 00:00:00	2005-09-12 12:00:00	1.220000
50%	8320.000000	2007-03-10 00:00:00	2009-03-19 00:00:00	2.780000
75%	11739.000000	2011-07-31 00:00:00	2013-09-11 12:00:00	4.750000
max	14971.000000	2014-11-11 00:00:00	2017-05-07 00:00:00	12.620000
std	4172.891992	NaN	NaN	2.397168

```
In [11]: # Summary Statistics for Categorical data:  
df.describe(exclude=[np.number])
```

	loan_type	repaid	loan_id	loan_start	loan_end
count	443	443.0	443.0	443	443
unique	4	2.0	443.0	NaN	NaN
top	home	1.0	10243.0	NaN	NaN
freq	121	237.0	1.0	NaN	NaN
mean	NaN	NaN	NaN	2007-08-02 12:56:53.092550912	2009-08-23 11:35:37.246049536
min	NaN	NaN	NaN	2000-01-26 00:00:00	2001-08-02 00:00:00
25%	NaN	NaN	NaN	2003-10-19 00:00:00	2005-09-12 12:00:00
50%	NaN	NaN	NaN	2007-03-10 00:00:00	2009-03-19 00:00:00
75%	NaN	NaN	NaN	2011-07-31 00:00:00	2013-09-11 12:00:00
max	NaN	NaN	NaN	2014-11-11 00:00:00	2017-05-07 00:00:00

6. Missing Values

```
In [12]: # use isnull().sum() to check for missing values  
df.isnull().sum()
```

```
Out[12]: loan_type      0  
loan_amount      0  
repaid          0  
loan_id          0  
loan_start       0  
loan_end         0  
rate             0  
dtype: int64
```


There are no missing values in the data.

Sk-learn library has an in-built function called Iterative Imputer to impute the missing values. Its sklearn documentation: <https://scikit-learn.org/stable/modules/generated/sklearn.impute.IterativeImputer.html>


7. Outliers Treatment

To check for the presence of outliers, we plot Boxplot.

```
In [13]: # For Loan_amount  
df['loan_amount'].plot(kind='box')  
plt.show()
```



```
In [14]: # For rate  
df['rate'].plot(kind='box')  
plt.show()
```



We can see that there are no outliers in the loan_amount column and some outliers are present in the rate column. To treat for outliers can either cap the values or transform the data. Shall demonstrate both the approaches here.

8. Transformation

8a. SQRT transformation

```
In [15]: df['SQRT_RATE'] = df['rate']**0.5
```

```
In [16]: df['sqrt_rate'] = np.sqrt(df['rate'])
```

```
In [17]: df.head()
```

```
Out[17]:   loan_type  loan_amount  repaid  loan_id  loan_start  loan_end  rate  SQRT_RATE  sqrt_rate  
client_id  
46109    home       13672      0   10243  2002-04-16  2003-12-20  2.15    1.466288  1.466288  
46109   credit       9794      0   10984  2003-10-21  2005-07-17  1.25    1.118034  1.118034  
46109    home       12734      1   10990  2006-02-01  2007-07-05  0.68    0.824621  0.824621  
46109   cash        12518      1   10596  2010-12-08  2013-05-05  1.24    1.113553  1.113553  
46109   credit       14049      1   11415  2010-07-07  2012-05-21  3.13    1.769181  1.769181
```

```
In [18]: #checking the skewness, kurtosis between the original and transformed data:
print("The skewness of the original data is {}".format(df.rate.skew()))
print('The skewness of the SQRT transformed data is {}'.format(df.SQRT_RATE.skew()))


print('')
print("The kurtosis of the original data is {}".format(df.rate.kurt()))
print("The kurtosis of the SQRT transformed data is {}".format(df.SQRT_RATE.kurt()))
```

The skewness of the original data is 0.884204614329943
The skewness of the SQRT transformed data is 0.04964154055528862
The kurtosis of the original data is 0.42437165143736433
The kurtosis of the SQRT transformed data is -0.6318437642052039

```
In [19]: # plotting the distribution

fig, axes = plt.subplots(1,2, figsize=(15,5))
sns.distplot(df['rate'], ax=axes[0])
sns.distplot(df['sqrt_rate'], ax=axes[1])

plt.show()
```



Result:

The Rate column was right skewed earlier. The skewness and kurtosis has reduced significantly. The transformed SQRT rate, on the right graph resembles normal distribution now.

8b. Log Transformation

```
In [20]: df['Log Rate'] = np.log(df['rate'])
```

```
In [21]: df.head()
```

```
Out[21]:   loan_type  loan_amount  repaid  loan_id  loan_start  loan_end  rate  SQRT_RATE  sqrt_rate  Log Rate
client_id
46109    home        13672      0  10243  2002-04-16  2003-12-20  2.15  1.466288  1.466288  0.765468
46109    credit       9794      0  10984  2003-10-21  2005-07-17  1.25  1.118034  1.118034  0.223144
46109    home        12734      1  10990  2006-02-01  2007-07-05  0.68  0.824621  0.824621  -0.385662
46109    cash         12518      1  10596  2010-12-08  2013-05-05  1.24  1.113553  1.113553  0.215111
46109    credit       14049      1  11415  2010-07-07  2012-05-21  3.13  1.769181  1.769181  1.141033
```

```
In [22]: print("The skewness of the original data is {}".format(df.rate.skew()))
print('The skewness of the SQRT transformed data is {}'.format(df.SQRT_RATE.skew()))
print("The skewness of the LOG transformed data is {}".format(df['Log Rate'].skew()))


print('')
print("The kurtosis of the original data is {}".format(df.rate.kurt()))
print("The kurtosis of the SQRT transformed data is {}".format(df.SQRT_RATE.kurt()))
print("The kurtosis of the LOG transformed data is {}".format(df['Log Rate'].kurt()))
```

The skewness of the original data is 0.884204614329943
The skewness of the SQRT transformed data is 0.04964154055528862
The skewness of the LOG transformed data is -1.5943217626331552
The kurtosis of the original data is 0.42437165143736433
The kurtosis of the SQRT transformed data is -0.6318437642052039
The kurtosis of the LOG transformed data is 4.157026150198228

```
In [23]: # plot the graph:
fig, axes = plt.subplots(1,3,figsize=(15,5))

sns.distplot(df['rate'], ax=axes[0])
sns.distplot(df['SQRT_RATE'], ax=axes[1])
sns.distplot(df['Log Rate'], ax=axes[2])

plt.show()
```



Inference:

Log Transformation made the rate left skewed and more peaked.

However, Log transformation is more closer to 0 and hence is more normal. Though it heavily manipulates the data.

In our case, square root transformation is more suitable.

```
In [24]: ## Using Lambda function :
df['LOG_Rate'] = df['rate'].apply(lambda x:np.log(x))
```

```
In [25]: df.head()
```

	loan_type	loan_amount	repaid	loan_id	loan_start	loan_end	rate	SQRT_RATE	sqrt_rate	Log Rate	LOG_Rate
client_id											
46109	home	13672	0	10243	2002-04-16	2003-12-20	2.15	1.466288	1.466288	0.765468	0.765468
46109	credit	9794	0	10984	2003-10-21	2005-07-17	1.25	1.118034	1.118034	0.223144	0.223144
46109	home	12734	1	10990	2006-02-01	2007-07-05	0.68	0.824621	0.824621	-0.385662	-0.385662
46109	cash	12518	1	10596	2010-12-08	2013-05-05	1.24	1.113553	1.113553	0.215111	0.215111
46109	credit	14049	1	11415	2010-07-07	2012-05-21	3.13	1.769181	1.769181	1.141033	1.141033

PROGRAM 3

Date: 12-04-2024

Use an appropriate data set for building the decision tree (ID3) and apply this knowledge to classify a new sample

Algorithm:

12/04/24

Program - 2 Decision Tree ID3

Algorithm:

ID3(Examples, Target-attribute, Attributes)

- Create a Root node for the tree.
- Return If all Examples are positive, Return the single-node tree Root, with Label = +
- If all Examples are negative, Return the single-node tree Root, with Label = -
- If attributes is empty, Return the single-node tree Root, with label = most common value of Target-attribute in Examples
- Otherwise Begin:
 - $A \leftarrow$ the attribute from Attributes that best* classifies Examples.
 - The decision attribute for Root $\leftarrow A$.
 - For each possible value, v_i , of A ,
 - Add a new tree branch below Root, corresponding to the test $A = v_i$
 - Let Examples _{v_i} be the subset of Examples that have values v_i for A .
 - If Examples _{v_i} is empty,
 - Then below this new branch add a leaf node with label = most common value of Target-attribute
 - Else below this new branch add the subtree ID3.
 - End
 - Return Root

Code:

Importing Database

```
In [2]: # Importing the required libraries
import pandas as pd
import numpy as np
import math

# Reading the dataset (Tennis-dataset)
data = pd.read_csv('/content/PlayTennis.csv')

In [ ]: from google.colab import drive
drive.mount('/content/drive')

In [3]: def highlight(cell_value):
    '''
    Highlight yes / no values in the dataframe
    '''
    color_1 = 'background-color: pink;'
    color_2 = 'background-color: lightgreen;'

    if cell_value == 'no':
        return color_1
    elif cell_value == 'yes':
        return color_2

    data.style.applymap(highlight)\n        .set_properties(subset=data.columns, **{'width': '100px'})\n        .set_table_styles([{'selector': 'th', 'props': [({'background-color': 'pink', 'color': 'white'}, ('border', '1px solid gray'), ('font-weight', 'bold')])},\n                          {'selector': 'tr:hover', 'props': [({'background-color': 'white'}, ('border', '1.5px solid black'))]}])\n\nOut[3]: outlook temp humidity windy play
0 sunny hot high False no
1 sunny hot high True no
2 overcast hot high False yes
3 rainy mild high False yes
4 rainy cool normal False yes
5 rainy cool normal True no
6 overcast cool normal True yes
7 sunny mild high False no
8 sunny cool normal False yes
9 rainy mild normal False yes
10 sunny mild normal True yes
11 overcast mild high True yes
12 overcast hot normal False yes
13 rainy mild high True no
```

Entropy of the dataset

```
In [4]:  
def find_entropy(data):  
    """  
    Returns the entropy of the class or features  
    formula: - Σ P(X)logP(X)  
    """  
    entropy = 0  
    for i in range(data.nunique()):  
        x = data.value_counts()[i]/data.shape[0]  
        entropy += (-x * math.log(x,2))  
    return round(entropy,3)  
  
def information_gain(data, data_):  
    """  
    Returns the information gain of the features  
    """  
    info = 0  
    for i in range(data_.nunique()):  
        df = data[data_ == data_.unique()[i]]  
        w_avg = df.shape[0]/data_.shape[0]  
        entropy = find_entropy(df.play)  
        x = w_avg * entropy  
        info += x  
    ig = find_entropy(data.play) - info  
    return round(ig, 3)  
  
def entropy_and_infogain(datax, feature):  
    """  
    Grouping features with the same class and computing their  
    entropy and information gain for splitting  
    """  
    for i in range(data[feature].nunique()):  
        df = datax[datax[feature]==data[feature].unique()[i]]  
        if df.shape[0] < 1:  
            continue  
  
        display(df[[feature, 'play']].style.applymap(highlight))  
        .set_properties(subset=[feature, 'play'], **{'width': '80px'})\  
        .set_table_styles([{'selector': 'th', 'props': [('background-color', 'lightgray'),  
                                                       ('border', '1px solid gray'),  
                                                       ('font-weight', 'bold')]}],  
                      {'selector': 'td', 'props': [('border', '1px solid gray')]},  
                      {'selector': 'tr:hover', 'props': [('background-color', 'white'),  
                                                       ('border', '1.5px solid black')]}])  
  
        print(f'Entropy of {feature} - {data[feature].unique()[i]} = {find_entropy(df.play)}')  
        print(f'Information Gain for {feature} = {information_gain(datax, datax[feature])}')
```

```
In [5]:  
print(f'Entropy of the entire dataset: {find_entropy(data.play)})'
```

Entropy of the entire dataset: 0.94

Entropy and Information Gain of temperature

```
In [6]: entropy_and_infogain(data, 'temp')
```

```
temp play
0 hot no
1 hot no
2 hot yes
12 hot yes
Entropy of temp - hot = 1.0
temp play
3 mild yes
7 mild no
9 mild yes
10 mild yes
11 mild yes
13 mild no
Entropy of temp - mild = 0.918
temp play
4 cool yes
5 cool no
6 cool yes
8 cool yes
Entropy of temp - cool = 0.811
Information Gain for temp = 0.029
```

Entropy and Information Gain of humidity

```
In [7]: entropy_and_infogain(data, 'humidity')
```

```
humidity play
0 high no
1 high no
2 high yes
3 high yes
7 high no
11 high yes
13 high no
Entropy of humidity - high = 0.985
humidity play
4 normal yes
5 normal no
6 normal yes
8 normal yes
9 normal yes
10 normal yes
12 normal yes
Entropy of humidity - normal = 0.592
Information Gain for humidity = 0.151
```

Entropy and Information Gain of windy

```
In [8]: entropy_and_infogain(data, 'windy')
```

```
windy play
```

```
0 False no
```

```
2 False yes
```

```
3 False yes
```

```
4 False yes
```

```
7 False no
```

```
8 False yes
```

```
9 False yes
```

```
12 False yes
```

```
Entropy of windy - False = 0.811
```

```
windy play
```

```
1 True no
```

```
5 True no
```

```
6 True yes
```

```
10 True yes
```

```
11 True yes
```

```
13 True no
```

```
Entropy of windy - True = 1.0
```

```
Information Gain for windy = 0.048
```

Rainy Outlook

Rainy -outlook

```
In [9]: rainy = data[data['outlook'] == 'rainy']
rainy.style.applymap(highlight)\n    .set_properties(subset=data.columns, **{'width': '100px'})\n    .set_table_styles([{'selector': 'th', 'props': [('background-color', 'lightgray'), ('border', '1px solid gray'),\n                                         ('font-weight', 'bold')]}],\n                     {'selector': 'tr:hover', 'props': [('background-color', 'white'), ('border', '1.5px solid black')]})
```

```
Out[9]:   outlook  temp  humidity  windy  play\n3      rainy   mild     high  False   yes\n4      rainy   cool    normal  False   yes\n5      rainy   cool    normal   True    no\n9      rainy   mild    normal  False   yes\n13     rainy   mild     high   True    no
```

```
In [10]: print(f'Entropy of the Rainy dataset: {find_entropy(rainy.play)}')
```

```
Entropy of the Rainy dataset: 0.971
```

```
In [11]: entropy_and_infogain(rainy, 'temp')
```

temp	play	
3	mild	yes
9	mild	yes
13	mild	no

Entropy of temp - mild = 0.918

temp	play	
4	cool	yes
5	cool	no

Entropy of temp - cool = 1.0
Information Gain for temp = 0.02

```
In [12]: entropy_and_infogain(rainy, 'humidity')
```

humidity	play	
3	high	yes
13	high	no

Entropy of humidity - high = 1.0

humidity	play	
4	normal	yes
5	normal	no
9	normal	yes

Entropy of humidity - normal = 0.918
Information Gain for humidity = 0.02

```
In [13]: entropy_and_infogain(rainy, 'windy')
```

windy	play	
3	False	yes
4	False	yes
9	False	yes

Entropy of windy - False = 0.0

windy	play	
5	True	no
13	True	no

Entropy of windy - True = 0.0
Information Gain for windy = 0.971

wind has highest information gain

Output

Output:

Entropy of the dataset : 0.9331

Pregnancies - Entropy : 3.482, IG : 0.062

Glucose - Entropy : 6.751, IG : 0.304

BloodPressure - Entropy : 4.792, IG : 0.059

SkinThickness - Entropy : 4.586, IG : 0.082

Insulin - Entropy : 4.682, IG : 0.277

BMI - Entropy : 7.594, IG : 0.344

DiabetesPedigreeFunction - Entropy : 8.829, IG : 0.65

Age - Entropy : 5.029, IG : 0.141

PROGRAM 4

Date: 19-04-2024

Implement Linear and Multi-Linear Regression algorithm using appropriate dataset

LINEAR REGRESSION:

Algorithm

03/05/2024

Program - 4

Implement Linear and Multi-Linear Regression algorithm.

Linear Regression:

```
function linear_regression (x, y, learning_rate, num_iterations)
    Initialize random values for slope (m) & intercept (b)
    for i = 1 to num_iterations :
        predictions = m * x + b
        errors = predictions - y
        loss = mean_squared_error(errors)
        gradient_m = (2/N) * sum(errors * x)
        gradient_b = (2/N) * sum(errors)
        m = m - learning_rate * gradient_m
        b = b - learning_rate * gradient_b
    Return m, b

function mean_squared_error (errors) :
    squared_errors = errors^2
    mse = sum(squared_errors) / sum(errors)
    return mse.
```

Code

Importing Dataset

```
In [1]: import pandas as pd  
import numpy as np  
import matplotlib.pyplot as plt
```

```
In [18]: df=pd.read_csv("/Salary_Data.csv")  
df
```

```
Out[18]:   YearsExperience    Salary  
0           1.1  39343.0  
1           1.3  46205.0  
2           1.5  37731.0  
3           2.0  43525.0  
4           2.2  39891.0  
5           2.9  56642.0  
6           3.0  60150.0  
7           3.2  54445.0  
8           3.2  64445.0  
9           3.7  57189.0  
10          3.9  63218.0  
11          4.0  55794.0  
12          4.0  56957.0  
13          4.1  57081.0  
14          4.5  61111.0  
15          4.9  67938.0  
16          5.1  66029.0  
17          5.3  83088.0  
18          5.9  81363.0  
19          6.0  93940.0  
20          6.8  91738.0  
21          7.1  98273.0  
22          7.9  101302.0  
23          8.2  113812.0  
24          8.7  109431.0  
25          9.0  105582.0  
26          9.5  116969.0  
27          9.6  112635.0  
28         10.3  122391.0  
29         10.5  121872.0
```

Slope and Intercept calculation

```
In [19]: def linear(X, b0, b1):  
    return [b0+b1*x for x in X]
```

```
In [20]: # b0 - Intercept  
def intercept(X, Y, b1):  
    x_ = np.mean(X)  
    y_ = np.mean(Y)  
  
    return y_-b1*x_
```

```
In [21]: # b1 - Slope  
def slope(X, Y):  
    x_ = np.mean(X)  
    y_ = np.mean(Y)  
  
    rise = sum([(x-x_) * (y-y_) for x,y in zip(X,Y)])  
    run = sum([(x-x_)**2 for x,y in zip(X,Y)])  
  
    return rise / run
```

Predicted Values Graph

```
In [23]: plt.figure(figsize=(8,5))  
plt.title("YearsExperience vs Salary")  
plt.scatter(predictor, target, color = "#4247ba")  
plt.xlabel("YearsExperience")  
plt.ylabel("Salary")  
plt.show()
```




```
In [24]: b1 = slope(predictor, target)  
b0 = intercept(predictor, target, b1)  
predicted = linear(predictor, b0, b1)
```

```
In [25]: plt.figure(figsize = (8, 5))  
plt.plot(predictor, predicted, color = "#f25f5c")  
plt.scatter(predictor, predicted, color = "#f25f5c") and Line is predicted data  
plt.title('Predicted values by Linear Regression', fontsize = 15)  
plt.xlabel("YearsExperience")  
plt.ylabel("Salary")  
plt.scatter(predictor, target, color = "#4247ba")  
plt.show()
```

Output

```
In [26]:  
print("Coefficients:\n=====")  
print("b0 : ", b0)  
print("b1 : ", b1)
```

```
Coefficients:  
=====  
b0 : 25792.20019866869  
b1 : 9449.962321455077
```



MULTIPLE LINEAR REGRESSION:

Algorithm

Multi-linear Regression

```
function initialize-parameters():
    randomly initialize  $\beta_0, \beta_1, \dots$ 
```

```
function hypothesis-function ( $x, \beta$ ):
     $h = \beta_0 + \beta_1 * x[1] + \beta_2 * x[2] + \dots + \beta_m * x[m]$ 
    return  $h$ 
```

```
function cost-function ( $x, y, \beta$ ):
    n = length( $x$ )
    total_error = 0
    for i = 1 to n:
        h = hypothesis-function ( $x[i], \beta$ )
        total_error += ( $h - y[i]$ )^2
    cost = (1/(2*n)) * total_error
    return cost.
```

```
function gradient-descent ( $x, y, \beta, \alpha, \text{iterations}, \text{threshold}$ ):
    n = length( $x$ )
    for i=0 to iterations:
        error_sum = 0
        for i = 1 to n:
            h = hypothesis-function ( $x[i], \beta$ )
            error =  $h - y[i]$ 
            error_sum += error
        for j = 0 to m:
             $\beta[j] = \beta[j] - \alpha * (1/n) * \text{error} * x[i][j]$ 
    cost = cost-function ( $x, y, \beta$ )
    if cost < threshold: break
    return  $\beta$ .
```

Code

```
In [19]: house = pd.read_csv('https://github.com/YBIFoundation/Dataset/raw/main/Boston.csv')

In [20]: house.head()

Out[20]:
   CRIM    ZN  INDUS  CHAS    NX    RM   AGE    DIS    RAD    TAX  PTRATIO     B   LSTAT    MEDV
0  0.00632  18.0    2.31     0  0.538  6.575  65.2  4.0900     1  296.0    15.3  396.90   4.98   24.0
1  0.02731    0.0    7.07     0  0.469  6.421  78.9  4.9671     2  242.0    17.8  396.90   9.14   21.6
2  0.02729    0.0    7.07     0  0.469  7.185  61.1  4.9671     2  242.0    17.8  392.83   4.03   34.7
3  0.03237    0.0    2.18     0  0.458  6.998  45.8  6.0622     3  222.0    18.7  394.63   2.94   33.4
4  0.06905    0.0    2.18     0  0.458  7.147  54.2  6.0622     3  222.0    18.7  396.90   5.33   36.2

In [21]: house.describe()

Out[21]:
      CRIM        ZN       INDUS      CHAS        NX        RM       AGE       DIS       RAD       TAX  PTRATI
count  506.000000  506.000000  506.000000  506.000000  506.000000  506.000000  506.000000  506.000000  506.000000  506.000000
mean   3.613524  11.363636  11.136779  0.069170  0.554695  6.284634  68.574901  3.795043  9.549407  408.237154  18.4555
std    8.601545  23.322453  6.860353  0.253994  0.115878  0.702617  28.148861  2.105710  8.707259  168.537116  2.1649
min    0.006320  0.000000  0.460000  0.000000  0.385000  3.561000  2.900000  1.129600  1.000000  187.000000  12.6000
25%   0.082045  0.000000  5.190000  0.000000  0.449000  5.885500  45.025000  2.100175  4.000000  279.000000  17.4000
50%   0.256510  0.000000  9.690000  0.000000  0.538000  6.208500  77.500000  3.207450  5.000000  330.000000  19.0500
75%   3.677083  12.500000  18.100000  0.000000  0.624000  6.623500  94.075000  5.188425  24.000000  666.000000  20.2000
max   88.976200 100.000000 27.740000  1.000000  0.871000  8.780000 100.000000 12.126500  24.000000  711.000000  22.0000

In [22]: house.columns

Out[22]: Index(['CRIM', 'ZN', 'INDUS', 'CHAS', 'NX', 'RM', 'AGE', 'DIS', 'RAD', 'TAX',
       'PTRATIO', 'B', 'LSTAT', 'MEDV'],
       dtype='object')

In [23]: y = house['MEDV']

In [24]: X = house.drop(['MEDV'],axis=1)

In [25]: from sklearn.model_selection import train_test_split
X_train, X_test, y_train, y_test = train_test_split(X,y, train_size=0.7, random_state=2529)

In [26]: X_train.shape, X_test.shape, y_train.shape, y_test.shape

Out[26]: ((350, 13), (150, 13), (350,), (150,))
```

```
In [25]: from sklearn.model_selection import train_test_split
X_train, X_test, y_train, y_test = train_test_split(X,y, train_size=0.7, random_state=2529)

In [26]: X_train.shape, X_test.shape, y_train.shape, y_test.shape

Out[26]: ((354, 13), (152, 13), (354,), (152,))

In [27]: from sklearn.linear_model import LinearRegression
model = LinearRegression()

In [28]: # Step 6 : train or fit model
model.fit(X_train,y_train)

Out[28]: LinearRegression()
In a Jupyter environment, please rerun this cell to show the HTML representation or trust the notebook.
On GitHub, the HTML representation is unable to render, please try loading this page with nbviewer.org.

In [29]: model.intercept_

Out[29]: 34.21916368862993

In [30]: model.coef_

Out[30]: array([-1.29e-01,  3.65e-02,  1.54e-02,  2.35e+00, -2.04e+01,  4.41e+00,
       4.61e-03, -1.59e+00,  2.51e-01, -9.60e-03, -9.64e-01,  1.01e-02,
      -5.43e-01])

In [31]: # Step 7 : predict model
y_pred = model.predict(X_test)
```

Output

```
In [32]: y_pred

Out[32]: array([31.72, 22.02, 21.17, 39.78, 20.1 , 22.86, 18.36, 14.79, 22.56,
       21.35, 18.38, 27.97, 29.86, 6.45, 10.68, 26.25, 21.89, 25.23,
       3.62, 36.22, 24.08, 22.94, 14.27, 20.79, 24.23, 16.74, 18.75,
       20.97, 28.51, 20.86, 9.23, 17.07, 22.07, 22.23, 39.26, 26.17,
       42.5 , 19.35, 34.52, 14.07, 13.81, 23.28, 11.79, 9.01, 21.65,
       25.55, 18.17, 16.82, 14.66, 14.86, 33.79, 33.27, 15.49, 24.08,
       27.64, 19.58, 45.02, 20.97, 20.07, 27.67, 34.59, 12.71, 23.66,
       31.66, 28.97, 32.46, 13.93, 35.49, 19.36, 19.6 , 1.44, 24.1 ,
       33.67, 20.62, 26.89, 21.29, 31.95, 29.74, 13.93, 13.82, 19.76,
       21.54, 20.87, 23.63, 28.8 , 23.64, 6.95, 22.2 , -6.82, 16.97,
       16.77, 25.44, 14.95, 3.72, 15.03, 16.91, 21.46, 31.66, 30.72,
       23.73, 22.19, 13.76, 18.47, 18.15, 36.6 , 27.49, 11. , 17.26,
       22.49, 16.53, 29.49, 22.89, 24.68, 20.38, 19.69, 22.55, 27.32,
       24.86, 20.2 , 29.14, 7.43, 5.85, 25.35, 38.73, 23.94, 25.28,
       20.11, 19.75, 25.07, 35.16, 27.32, 27.26, 31.4 , 16.55, 14.3 ,
       23.77, 7.65, 23.35, 21.37, 26.12, 25.32, 13.12, 17.67, 36.2 ,
       20.5 , 27.95, 22.46, 18.15, 31.24, 20.85, 27.36, 30.53])
```

```
In [33]: # Step 8 : model accuracy
from sklearn.metrics import mean_absolute_error, mean_absolute_percentage_error, mean_squared_error
```

```
In [34]: mean_absolute_error(y_test,y_pred)
```

```
Out[34]: 3.155030927602485
```

PROGRAM 5

Date:03-05-2024

Build Logistic Regression Model for a given dataset

Algorithm

03/05/24
Program-5

Build logistic Regression Model for a given dataset.

```
function logistic-regression(x, y, learning-rate, num-itr):
    Initialize random values for weights (w) & bias (b)
    for i = 1 to num-itr:
        logits = x * w + b
        pred = sigmoid(logits)
        loss = compute-loss(y, pred)
        update weights & bias using gradients
    return w, b

function sigmoid(x):
    return 1 / (1 + exp(-x))

function compute-loss(y-true, y-pred):
    loss = mean(y-true * log(y-pred) + (1-y-true) * log(1-y-pred))
    return loss
```

Code

```
In [2]: import numpy as np # linear algebra
import pandas as pd # data processing, CSV file I/O (e.g. pd.read_csv)
import matplotlib.pyplot as plt
# Input data files are available in the "../input/" directory.
# For example, running this (by clicking run or pressing Shift+Enter) will list the files in the input directory

import os

In [4]: data = pd.read_csv('/content/data.csv')

In [5]: data.drop(['Unnamed: 32','id'], axis=1, inplace=True)
data.diagnosis = [1 if each == "M" else 0 for each in data.diagnosis]
y = data.diagnosis.values
x_data = data.drop(['diagnosis'], axis=1)

In [7]: # Assuming x_data is a numpy array or pandas DataFrame
x = (x_data - np.min(x_data)) / (np.max(x_data) - np.min(x_data))

In [8]: from sklearn.model_selection import train_test_split

x_train, x_test, y_train, y_test = train_test_split(x, y, test_size=0.15, random_state=42)

x_train = x_train.T
x_test = x_test.T
y_train = y_train.T
y_test = y_test.T

print("x train: ",x_train.shape)
print("x test: ",x_test.shape)
print("y train: ",y_train.shape)
print("y test: ",y_test.shape)

x train: (30, 483)
x test: (30, 86)
y train: (483,)
y test: (86,)

In [9]: def initialize_weights_and_bias(dimension):
    w = np.full((dimension,1),0.01)
    b = 0.0
    return w, b

In [10]: def sigmoid(z):
    y_head = 1/(1+np.exp(-z))
    return y_head

In [ ]: def forward_backward_propagation(w,b,x_train,y_train):
    # forward propagation
    z = np.dot(w.T,x_train) + b
    y_head = sigmoid(z)
    loss = -y_train*np.log(y_head)-(1-y_train)*np.log(1-y_head)
    cost = (np.sum(loss))/x_train.shape[1]      # x_train.shape[1] is for scaling
    # backward propagation
    derivative_weight = (np.dot(x_train,((y_head-y_train).T)))/x_train.shape[1] # x_train.shape[1] is for scaling
    derivative_bias = np.sum(y_head-y_train)/x_train.shape[1]                  # x_train.shape[1] is for scaling
    gradients = {"derivative_weight": derivative_weight,"derivative_bias": derivative_bias}
    return cost,gradients
```

```
In [ ]: def update(w, b, x_train, y_train, learning_rate, number_of_iterarion):
    cost_list = []
    cost_list2 = []
    index = []
    # updating(learning) parameters is number_of_iterarion times
    for i in range(number_of_iterarion):
        # make forward and backward propagation and find cost and gradients
        cost,gradients = forward_backward_propagation(w,b,x_train,y_train)
        cost_list.append(cost)
        w = w - learning_rate * gradients["derivative_weight"]
        b = b - learning_rate * gradients["derivative_bias"]
        if i % 10 == 0:
            cost_list2.append(cost)
            index.append(i)
            print ("Cost after iteration %i: %f" %(i, cost))
    # we update(Learn) parameters weights and bias
    parameters = {"weight": w,"bias": b}
    plt.plot(index,cost_list2)
    plt.xticks(index,rotation='vertical')
    plt.xlabel("Number of Iterarion")
    plt.ylabel("Cost")
    plt.show()
    return parameters, gradients, cost_list
```

```
In [11]: def predict(w,b,x_test):
    # x_test is a input for forward propagation
    z = sigmoid(np.dot(w.T,x_test)+b)
    Y_prediction = np.zeros((1,x_test.shape[1]))
    # if z is bigger than 0.5, our prediction is sign one (y_head=1),
    # if z is smaller than 0.5, our prediction is sign zero (y_head=0),
    for i in range(z.shape[1]):
        if z[0,i]<= 0.5:
            Y_prediction[0,i] = 0
        else:
            Y_prediction[0,i] = 1
    return Y_prediction
```

```
In [18]: def sigmoid(z):
    return 1 / (1 + np.exp(-z))

def initialize_weights_and_bias(dim):
    w = np.zeros((dim, 1))
    b = 0
    return w, b

def compute_cost(w, b, x, y):
    m = x.shape[1]
    A = sigmoid(np.dot(w.T, x) + b)
    cost = -1 / m * np.sum(y * np.log(A) + (1 - y) * np.log(1 - A))
    return cost

def propagate(w, b, x, y):
    m = x.shape[1]
    A = sigmoid(np.dot(w.T, x) + b)
    dw = 1 / m * np.dot(x, (A - y).T)
    db = 1 / m * np.sum(A - y)
    return dw, db
```

```

def logistic_regression(x_train, y_train, x_test, y_test, learning_rate, num_iterations):
    # Initialize
    dimension = x_train.shape[0] # Number of features
    w, b = initialize_weights_and_bias(dimension)
    costs = []

    # Gradient Descent
    for i in range(num_iterations):
        # Forward and Backward Propagation
        dw, db = propagate(w, b, x_train, y_train)

        # Update parameters
        w -= learning_rate * dw
        b -= learning_rate * db

        # Record the costs
        if i % 100 == 0:
            cost = compute_cost(w, b, x_train, y_train)
            costs.append(cost)
            print(f"Cost after iteration {i}: {cost}")

    # Evaluate model
    y_prediction_train = predict(w, b, x_train)
    y_prediction_test = predict(w, b, x_test)

    train_accuracy = 100 - np.mean(np.abs(y_prediction_train - y_train)) * 100
    test_accuracy = 100 - np.mean(np.abs(y_prediction_test - y_test)) * 100

    print("Train accuracy: {}".format(train_accuracy))
    print("Test accuracy: {}".format(test_accuracy))

    return w, b

# Assuming you have defined the predict function
# def predict(w, b, x):
#     ...

# Assuming you have defined x_train, y_train, x_test, y_test, learning_rate, and num_iterations
logistic_regression(x_train, y_train, x_test, y_test, learning_rate=1, num_iterations=100)

```

Output

```
Cost after iteration 0: 0.6782740160052536
Train accuracy: 80.74534161490683 %
Test accuracy: 81.3953488372093 %


Out[18]: (array([[ 1.77806654e-02],
       [ 1.10160388e-02],
       [ 1.27806976e-01],
       [ 1.95749649e+00],
       [ 1.85931875e-05],
       [ 2.68863405e-04],
       [ 4.89020048e-04],
       [ 2.63106803e-04],
       [ 3.49357933e-05],
       [-2.02145931e-05],
       [ 1.25690784e-03],
       [-3.98285024e-04],
       [ 8.96937014e-03],
       [ 2.02426962e-01],
       [-3.60718647e-06],
       [ 4.19150446e-05],
       [ 6.03411729e-05],
       [ 2.00740406e-05],
       [-6.24803672e-06],
       [ 6.24944780e-07],
       [ 2.79506973e-02],
       [ 1.99326360e-02],
       [ 1.98774929e-01],
       [ 3.39189908e+00],
       [ 5.79135019e-05],
       [ 8.53041205e-04],
       [ 1.25862280e-03],
       [ 4.60695564e-04],
       [ 1.89671301e-04],
       [ 3.52490835e-05]]),
 -1.5161875221606185)
```

PROGRAM 6

Date: 19-04-2024

Build KNN Classification model for a given dataset.

Algorithm



Code

```
In [1]: import numpy as np # linear algebra
import pandas as pd # data processing, CSV file I/O (e.g. pd.read_csv)
import matplotlib.pyplot as plt # for data visualization purposes
import seaborn as sns # for data visualization
%matplotlib inline

In [2]: data = '/content/cancer_detector.txt'
df = pd.read_csv(data, header=None)

In [3]: df.shape

Out[3]: (699, 11)

In [4]: col_names = ['Id', 'Clump_thickness', 'Uniformity_Cell_Size', 'Uniformity_Cell_Shape', 'Marginal_Adhesion',
   'Single_Epithelial_Cell_Size', 'Bare_Nuclei', 'Bland_Chromatin', 'Normal_Nucleoli', 'Mitoses', 'Class']

df.columns = col_names

df.columns

Out[4]: Index(['Id', 'Clump_thickness', 'Uniformity_Cell_Size',
   'Uniformity_Cell_Shape', 'Marginal_Adhesion',
   'Single_Epithelial_Cell_Size', 'Bare_Nuclei', 'Bland_Chromatin',
   'Normal_Nucleoli', 'Mitoses', 'Class'],
   dtype='object')

In [5]: df.head()

Out[5]:    Id Clump_thickness Uniformity_Cell_Size Uniformity_Cell_Shape Marginal_Adhesion Single_Epithelial_Cell_Size Bare_Nuclei Bla
  0  1000025          5              1                  1                  1                      2                  1
  1  1002945          5              4                  4                  5                      7                 10
  2  1015425          3              1                  1                  1                      2                  2
  3  1016277          6              8                  8                  1                      3                  4
  4  1017023          4              1                  1                  3                      2                  1
```

```
In [18]: import numpy as np
```

```
In [21]: # view summary statistics in numerical variables
print(round(df.describe(),2))

   Clump_Thickness Uniformity_Cell_Size Uniformity_Cell_Shape \
count          699.00            699.00            699.00
mean           4.42             3.13             3.21
std            2.82             3.05             2.97
min            1.00             1.00             1.00
25%            2.00             1.00             1.00
50%            4.00             1.00             1.00
75%            6.00             5.00             5.00
max           10.00            10.00            10.00

   Marginal_Adhesion Single_Epithelial_Cell_Size Bare_Nuclei \
count          699.00            699.00            683.00
mean           2.81             3.22             3.54
std            2.86             2.21             3.64
min            1.00             1.00             1.00
25%            1.00             2.00             1.00
50%            1.00             2.00             1.00
75%            4.00             4.00             6.00
max           10.00            10.00            10.00

   Bland_Chromatin Normal_Nucleoli Mitoses Class
count          699.00            699.00            699.00
mean           3.44             2.87             1.59             2.69
std            2.44             3.05             1.72             0.95
min            1.00             1.00             1.00             2.00
25%            2.00             1.00             1.00             2.00
50%            3.00             1.00             1.00             2.00
75%            5.00             4.00             1.00             4.00
max           10.00            10.00            10.00            4.00

In [22]: X = df.drop(['Class'], axis=1)
y = df['Class']

In [23]: from sklearn.model_selection import train_test_split
X_train, X_test, y_train, y_test = train_test_split(X, y, test_size = 0.2, random_state = 0)

In [24]: X_train.shape, X_test.shape
Out[24]: ((559, 9), (140, 9))

In [25]: for col in X_train.columns:
    if X_train[col].isnull().mean()>0:
        print(col, round(X_train[col].isnull().mean(),4))
Bare_Nuclei 0.0233

In [26]: for df1 in [X_train, X_test]:
    for col in X_train.columns:
        col_median=X_train[col].median()
        df1[col].fillna(col_median, inplace=True)

In [27]: cols = X_train.columns

In [28]: from sklearn.preprocessing import StandardScaler
scaler = StandardScaler()

X_train = scaler.fit_transform(X_train)

X_test = scaler.transform(X_test)

In [29]: X_train = pd.DataFrame(X_train, columns=[cols])

In [30]: X_test = pd.DataFrame(X_test, columns=[cols])
```

Output

```
In [34]: knn.predict_proba(X_test)[:,0]
```

```
Out[34]: array([1.          , 1.          , 0.33333333, 1.          , 0.          ,
   1.          , 0.          , 1.          , 0.          , 0.66666667,
   1.          , 1.          , 0.          , 0.33333333, 0.          ,
   1.          , 1.          , 0.          , 0.          , 1.          ,
   0.          , 0.          , 1.          , 1.          , 1.          ,
   0.          , 1.          , 1.          , 0.          , 0.          ,
   1.          , 1.          , 1.          , 1.          , 1.          ,
   0.66666667, 1.          , 0.          , 1.          , 1.          ,
   1.          , 1.          , 1.          , 1.          , 0.          ,
   0.          , 1.          , 0.          , 1.          , 0.          ,
   0.          , 1.          , 1.          , 0.          , 1.          ,
   1.          , 1.          , 1.          , 0.66666667, 1.          ,
   0.          , 1.          , 1.          , 0.          , 0.          ,
   0.33333333, 0.          , 1.          , 1.          , 0.          ,
   1.          , 1.          , 0.          , 0.          , 1.          ,
   1.          , 1.          , 1.          , 0.          , 1.          ,
   1.          , 1.          , 0.          , 1.          , 1.          ,
   1.          , 0.          , 1.          , 0.          , 0.          ,
   1.          , 1.          , 0.66666667, 0.          , 1.          ,
   1.          , 1.          , 0.          , 1.          , 0.          ,
   0.          , 1.          , 1.          , 1.          , 0.          ,
   1.          , 1.          , 1.          , 1.          , 1.          ,
   0.          , 0.33333333, 0.          , 1.          , 1.          ,
   1.          , 1.          , 1.          , 0.          , 0.          ,
   0.          , 0.33333333, 1.          , 0.          , 1.          ,
   1.          , 0.33333333, 0.33333333, 0.          , 0.          ,
   0.          , 1.          , 1.          , 0.33333333, 0.          ,
   1.          , 1.          , 0.          , 1.          , 1.          ,
   ])
```

```
In [35]: from sklearn.metrics import accuracy_score  
  
print('Model accuracy score: {:.4f}'.format(accuracy_score(y_test, y_pred)))  
  
Model accuracy score: 0.9714
```

```
In [36]: y_pred_train = knn.predict(X_train)
```

```
In [37]: print('Training-set accuracy score: {:.4f}'.format(accuracy_score(y_train, y_pred_train)))
```


Training-set accuracy score: 0.9821

PROGRAM 7

Date:24-05-2024

Build Support vector machine model for a given dataset

Algorithm



Code

```
Open In Colab
```

```
In [ ]: import seaborn as sns
import numpy as np
import pandas as pd
import matplotlib.pyplot as plt
import plotly.express as px
```

```
In [ ]: df = pd.read_csv('/content/breast-cancer.csv')
df.head()
```

```
Out[ ]:
```

	id	diagnosis	radius_mean	texture_mean	perimeter_mean	area_mean	smoothness_mean	compactness_mean	concavity_mean
0	842302	M	17.99	10.38	122.80	1001.0	0.11840	0.27760	0.300
1	842517	M	20.57	17.77	132.90	1326.0	0.08474	0.07954	0.086
2	84300903	M	19.69	21.25	130.00	1208.0	0.10960	0.15990	0.197
3	84348801	M	11.42	20.38	77.58	386.1	0.14250	0.28390	0.241
4	84358402	M	20.29	14.34	135.10	1297.0	0.10030	0.13280	0.198

5 rows × 32 columns

```
In [ ]: df.drop('id', axis=1, inplace=True) #drop redundant columns
```

```
In [ ]: df.describe().T
```

```
Out[ ]:
```

	count	mean	std	min	25%	50%	75%	max
radius_mean	569.0	14.127292	3.524049	6.981000	11.700000	13.370000	15.780000	28.11000
texture_mean	569.0	19.289649	4.301036	9.710000	16.170000	18.840000	21.800000	39.28000
perimeter_mean	569.0	91.969033	24.298981	43.790000	75.170000	86.240000	104.10000	188.50000
area_mean	569.0	654.889104	351.914129	143.500000	420.300000	551.100000	782.700000	2501.00000
smoothness_mean	569.0	0.096360	0.014064	0.052630	0.086370	0.095870	0.105300	0.16340
compactness_mean	569.0	0.104341	0.052813	0.019980	0.064920	0.092630	0.130400	0.34540
concavity_mean	569.0	0.088799	0.079720	0.000000	0.029560	0.061540	0.130700	0.42680
concave points_mean	569.0	0.048919	0.088803	0.000000	0.020310	0.033500	0.074000	0.20120
symmetry_mean	569.0	0.181162	0.027414	0.106000	0.161900	0.179200	0.195700	0.30400
fractal_dimension_mean	569.0	0.062798	0.007060	0.049960	0.057700	0.061540	0.066120	0.09744
radius_se	569.0	0.405172	0.277313	0.111500	0.232400	0.324200	0.478900	2.87300
texture_se	569.0	1.216853	0.551648	0.360200	0.833900	1.108000	1.474000	4.88500
perimeter_se	569.0	2.066059	2.021855	0.757000	1.606000	2.287000	3.357000	21.98000
area_se	569.0	40.337079	45.491006	6.802000	17.850000	24.530000	45.190000	542.20000
smoothness_se	569.0	0.007041	0.003003	0.001713	0.005169	0.006380	0.008146	0.09113
compactness_se	569.0	0.025478	0.017908	0.002252	0.013080	0.020450	0.032450	0.13540
concavity_se	569.0	0.031894	0.080186	0.000000	0.015090	0.025890	0.042050	0.39600
concave points_se	569.0	0.011796	0.006170	0.000000	0.007688	0.010930	0.014710	0.05279
symmetry_se	569.0	0.020542	0.008266	0.007882	0.015160	0.018730	0.023480	0.07895
fractal_dimension_se	569.0	0.005795	0.002646	0.000095	0.002248	0.003187	0.004558	0.02964
radius_worst	569.0	16.269190	4.833242	7.980000	13.010000	14.970000	18.790000	36.04000
texture_worst	569.0	25.677223	6.146258	12.020000	21.080000	25.410000	29.720000	49.54000
perimeter_worst	569.0	107.261213	33.602542	50.410000	84.110000	97.660000	125.400000	251.20000
area_worst	569.0	880.583128	569.356993	185.200000	515.300000	686.500000	1084.000000	4254.00000
smoothness_worst	569.0	0.132369	0.022832	0.071170	0.116600	0.131300	0.146000	0.22260
compactness_worst	569.0	0.254265	0.157336	0.027290	0.147200	0.211900	0.339100	1.05800
concavity_worst	569.0	0.272188	0.208624	0.000000	0.114500	0.126700	0.382900	1.25200
concave points_worst	569.0	0.114606	0.065732	0.000000	0.064980	0.099930	0.161400	0.29100
symmetry_worst	569.0	0.290076	0.061867	0.156500	0.250400	0.282200	0.317900	0.66380
fractal_dimension_worst	569.0	0.083946	0.018061	0.055040	0.071460	0.080040	0.092080	0.20750

```

In [ ]: df['diagnosis'] = (df['diagnosis'] == 'M').astype(int) #encode the label into 1/0

In [ ]: corr = df.corr()

In [ ]: # Get the absolute value of the correlation
cor_target = abs(corr['diagnosis'])

# Select Highly correlated features (threshold = 0.2)
relevant_features = cor_target[cor_target>0.2]

# Collect the names of the features
names = [index for index, value in relevant_features.items()]

# Drop the target variable from the results
names.remove('diagnosis')

# Display the results
print(names)

['radius_mean', 'texture_mean', 'perimeter_mean', 'area_mean', 'smoothness_mean', 'compactness_mean', 'concave points_mean', 'symmetry_mean', 'radius_se', 'perimeter_se', 'area_se', 'compactness_se', 'concavity_se', 'concave points_se', 'radius_worst', 'texture_worst', 'perimeter_worst', 'area_worst', 'smoothness_worst', 'compactness_worst', 'concavity_worst', 'concave points_worst', 'symmetry_worst', 'fractal_dimension_worst']

In [ ]: X = df[names].values
y = df['diagnosis']

In [ ]: def scale(X):
    """
    Standardizes the data in the array X.

    Parameters:
    X (numpy.ndarray): Features array of shape (n_samples, n_features).

    Returns:
    numpy.ndarray: The standardized features array.
    """
    # Calculate the mean and standard deviation of each feature
    mean = np.mean(X, axis=0)
    std = np.std(X, axis=0)

    # Standardize the data
    X = (X - mean) / std
    return X

In [ ]: X = scale(X)

In [ ]: def train_test_split(X, y, random_state=42, test_size=0.2):
    """
    Splits the data into training and testing sets.

    Parameters:
    X (numpy.ndarray): Features array of shape (n_samples, n_features).
    y (numpy.ndarray): Target array of shape (n_samples,).
    random_state (int): Seed for the random number generator. Default is 42.
    test_size (float): Proportion of samples to include in the test set. Default is 0.2.

    Returns:
    Tuple[numpy.ndarray]: A tuple containing X_train, X_test, y_train, y_test.
    """
    # Get number of samples
    n_samples = X.shape[0]

    # Set the seed for the random number generator
    np.random.seed(random_state)

    # Shuffle the indices
    shuffled_indices = np.random.permutation(np.arange(n_samples))

    # Determine the size of the test set
    test_size = int(n_samples * test_size)

    # Split the indices into test and train
    test_indices = shuffled_indices[:test_size]
    train_indices = shuffled_indices[test_size:]

    # Split the features and target arrays into test and train
    X_train, X_test = X[train_indices], X[test_indices]
    y_train, y_test = y[train_indices], y[test_indices]

    return X_train, X_test, y_train, y_test

In [ ]: X_train, X_test, y_train, y_test = train_test_split(X, y, test_size = 0.2, random_state=42) #split the data into training and testing sets

```

```
In [ ]: class SVM:

    def __init__(self, iterations=1000, lr=0.01, lambdaa=0.01):
        self.lambdaa = lambdaa
        self.iterations = iterations
        self.lr = lr
        self.w = None
        self.b = None
    def initialize_parameters(self, X):
        m, n = X.shape
        self.w = np.zeros(n)
        self.b = 0

    def gradient_descent(self, X, y):
        y_ = np.where(y <= 0, -1, 1)
        for i, x in enumerate(X):
            if y_[i] * (np.dot(x, self.w) - self.b) >= 1:
                dw = 2 * self.lambdaa * self.w
                db = 0
            else:
                dw = 2 * self.lambdaa * self.w - np.dot(x, y_[i])
                db = y_[i]
            self.update_parameters(dw, db)

    def update_parameters(self, dw, db):
        self.w = self.w - self.lr * dw
        self.b = self.b - self.lr * db
    def fit(self, X, y):
        self.initialize_parameters(X)
        for i in range(self.iterations):
            self.gradient_descent(X, y)

    def predict(self, X):
        # get the outputs
        output = np.dot(X, self.w) - self.b
        # get the signs of the labels depending on if it's greater/less than zero
        label_signs = np.sign(output)
        # set predictions to 0 if they are less than or equal to -1 else set them to 1
        predictions = np.where(label_signs <= -1, 0, 1)
        return predictions
```

```
In [ ]: def accuracy(y_true, y_pred):

    total_samples = len(y_true)
    correct_predictions = np.sum(y_true == y_pred)
    return (correct_predictions / total_samples)
```

Output

```
In [ ]: model = SVM()
model.fit(X_train,y_train)
predictions = model.predict(X_test)

accuracy(y_test, predictions)
```

```
Out[ ]: 0.9823008849557522
```

PROGRAM 8

Date: 31-05-2024

Build Artificial Neural Network model with back propagation on a given dataset

Algorithm

31/5/24 Date _____
Page _____

Program 8

Build Artificial Neural Network model with back propagation on a given dataset.

Algorithm :

1. Create a feed-forward network with n_i inputs, n_{hidden} hidden units, and n_{out} output units.
2. Initialize all network weights to small random numbers.
3. Until the termination condition is met, do
 - For each (\bar{x}, \bar{E}) , in training examples, do
 - Propagate the input forward through the network:
 - Input the instance \bar{x} , to the network and compute the output o_u of every unit u in the network.
 - Propagate the error backward through the network:
 - For each network output unit k , calculate its error term δ_k
 $\delta_k \leftarrow o_k(1-o_k)(t_k - o_k)$
 - For each hidden unit h , calculate its error term δ_h
 $\delta_h \leftarrow o_h(1-o_h) \sum_{k \in \text{outputs}} w_{hk} \delta_k$
 - Update each network weight w_{ji}
 $w_{ji} \leftarrow w_{ji} + \Delta w_{ji}$
where $\Delta w_{ji} = \eta \delta_j x_{ij}$

Output : Learning Accuracy = 0.66233, Testing Accuracy = 1.0

Code

```
In [1]: import numpy as np
from sklearn.model_selection import train_test_split

db = np.loadtxt("/content/duke-breast-cancer.txt")
print("Database raw shape (%s,%s)" % np.shape(db))

Database raw shape (86,7130)

In [2]: np.random.shuffle(db)
y = db[:, 0]
x = np.delete(db, [0], axis=1)
x_train, x_test, y_train, y_test = train_test_split(x, y, test_size=0.1)
print(np.shape(x_train),np.shape(x_test))

(77, 7129) (9, 7129)

In [3]: hidden_layer = np.zeros(72)
weights = np.random.random((len(x[0]), 72))
output_layer = np.zeros(2)
hidden_weights = np.random.random((72, 2))

In [4]: def sum_function(weights, index_locked_col, x):
    result = 0
    for i in range(0, len(x)):
        result += x[i] * weights[i][index_locked_col]
    return result

In [5]: def activate_layer(layer, weights, x):
    for i in range(0, len(layer)):
        layer[i] = 1.7159 * np.tanh(2.0 * sum_function(weights, i, x) / 3.0)

In [6]: def soft_max(layer):
    soft_max_output_layer = np.zeros(len(layer))
    for i in range(0, len(layer)):
        denominator = 0
        for j in range(0, len(layer)):
            denominator += np.exp(layer[j] - np.max(layer))
        soft_max_output_layer[i] = np.exp(layer[i] - np.max(layer)) / denominator
    return soft_max_output_layer

In [7]: def recalculate_weights(learning_rate, weights, gradient, activation):
    for i in range(0, len(weights)):
        for j in range(0, len(weights[i])):
            weights[i][j] = (learning_rate * gradient[j] * activation[i]) + weights[i][j]

In [8]: def back_propagation(hidden_layer, output_layer, one_hot_encoding, learning_rate, x):
    output_derivative = np.zeros(2)
    output_gradient = np.zeros(2)
    for i in range(0, len(output_layer)):
        output_derivative[i] = (1.0 - output_layer[i]) * output_layer[i]
    for i in range(0, len(output_layer)):
        output_gradient[i] = output_derivative[i] * (one_hot_encoding[i] - output_layer[i])
    hidden_derivative = np.zeros(72)
    hidden_gradient = np.zeros(72)
    for i in range(0, len(hidden_layer)):
        hidden_derivative[i] = (1.0 - hidden_layer[i]) * (1.0 + hidden_layer[i])
    for i in range(0, len(hidden_layer)):
        sum_ = 0
        for j in range(0, len(output_gradient)):
            sum_ += output_gradient[j] * hidden_weights[i][j]
        hidden_gradient[i] = sum_ * hidden_derivative[i]
    recalculate_weights(learning_rate, hidden_weights, output_gradient, hidden_layer)
    recalculate_weights(learning_rate, weights, hidden_gradient, x)
```

Output

```
In [9]: one_hot_encoding = np.zeros((2,2))
for i in range(0, len(one_hot_encoding)):
    one_hot_encoding[i][i] = 1
training_correct_answers = 0
for i in range(0, len(x_train)):
    activate_layer(hidden_layer, weights, x_train[i])
    activate_layer(output_layer, hidden_weights, hidden_layer)
    output_layer = soft_max(output_layer)
    training_correct_answers += 1 if y_train[i] == np.argmax(output_layer) else 0
    back_propagation(hidden_layer, output_layer, one_hot_encoding[int(y_train[i])], -1, x_train[i])
print("MLP Correct answers while learning: %s / %s (Accuracy = %s) on %s database." % (training_correct_answers, len(x_train),
                                                                                     training_correct_answers/len(x_train), "Duke breast cancer database"))

MLP Correct answers while learning: 51 / 77 (Accuracy = 0.6623376623376623) on Duke breast cancer database.
```

```
In [10]: testing_correct_answers = 0
for i in range(0, len(x_test)):
    activate_layer(hidden_layer, weights, x_test[i])
    activate_layer(output_layer, hidden_weights, hidden_layer)
    output_layer = soft_max(output_layer)
    testing_correct_answers += 1 if y_test[i] == np.argmax(output_layer) else 0
print("MLP Correct answers while testing: %s / %s (Accuracy = %s) on %s database" % (testing_correct_answers, len(x_test),
                                                                                     testing_correct_answers/len(x_test), "Duke breast cancer database"))


MLP Correct answers while testing: 9 / 9 (Accuracy = 1.0) on Duke breast cancer database
```

PROGRAM 9

Date: 31-05-2024

a) Implement Random forest ensemble method on a given dataset.

Algorithm



Code

[Open in Colab](#)

```
In [1]: # Pandas is used for data manipulation
import pandas as pd
# Read in data and display first 5 rows
features = pd.read_csv('/content/temp.csv')
features.head(5)
```

```
Out[1]:   year month day week temp_2 temp_1 average actual forecast_noaa forecast_acc forecast_under friend
0 2016    1    1   Fri     45     45   45.6    45        43        50        44      29
1 2016    1    2   Sat     44     45   45.7    44        41        50        44      61
2 2016    1    3   Sun     45     44   45.8    41        43        46        47      56
3 2016    1    4   Mon     44     41   45.9    40        44        48        46      53
4 2016    1    5  Tues     41     40   46.0    44        46        46        46      41
```

```
In [2]: print('The shape of our features is:', features.shape)
```

The shape of our features is: (348, 12)

```
In [3]: # Descriptive statistics for each column
features.describe()
```

```
Out[3]:   year   month   day   temp_2   temp_1   average   actual   forecast_noaa   forecast_acc   forecast_under
count  348.0  348.000000  348.000000  348.000000  348.000000  348.000000  348.000000  348.000000  348.000000  348.000000
mean  2016.0  6.477011  15.514368  62.652299  62.701149  59.760632  62.543103  57.238506  62.373563  59.772989  60.0
std   0.0  3.498380  8.772982  12.165398  12.120542  10.527306  11.794146  10.605746  10.549381  10.705256  15.0
min   2016.0  1.000000  1.000000  35.000000  35.000000  45.100000  35.000000  41.000000  46.000000  44.000000  28.0
25%  2016.0  3.000000  8.000000  54.000000  54.000000  49.975000  54.000000  48.000000  53.000000  50.000000  47.0
50%  2016.0  6.000000  15.000000  62.500000  62.500000  58.200000  62.500000  56.000000  61.000000  58.000000  60.0
75%  2016.0  10.000000  23.000000  71.000000  71.000000  69.025000  71.000000  66.000000  72.000000  69.000000  71.0
max  2016.0  12.000000  31.000000  117.000000  117.000000  77.400000  92.000000  77.000000  82.000000  79.000000  95.0
```

```
In [4]: # One-hot encode the data using pandas get_dummies
features = pd.get_dummies(features)
# Display the first 5 rows of the last 12 columns
features.iloc[:,5:].head(5)
```

```
Out[4]:   average   actual   forecast_noaa   forecast_acc   forecast_under   friend   week_Fri   week_Mon   week_Sat   week_Sun   week_Thurs   week_Tu
0  45.6      45          43          50          44      29      True     False     False     False     False     False
1  45.7      44          41          50          44      61     False     False     True     False     False     False
2  45.8      41          43          46          47      56     False     False     False     True     False     False
3  45.9      40          44          48          46      53     False     True     False     False     False     False
4  46.0      44          46          46          46      41     False     False     False     False     False     True
```

```
In [5]: # Use numpy to convert to arrays
import numpy as np
# Labels are the values we want to predict
labels = np.array(features['actual'])
# Remove the labels from the features
# axis 1 refers to the columns
features= features.drop('actual', axis = 1)
# Saving feature names for later use
feature_list = list(features.columns)
# Convert to numpy array
features = np.array(features)
```

```
In [6]: # Using Skicit-learn to split data into training and testing sets
from sklearn.model_selection import train_test_split
# Split the data into training and testing sets
train_features, test_features, train_labels, test_labels = train_test_split(features, labels, test_size = 0.25, random_state = 42)
```

```
In [7]: print('Training Features Shape:', train_features.shape)
print('Training Labels Shape:', train_labels.shape)
print('Testing Features Shape:', test_features.shape)
print('Testing Labels Shape:', test_labels.shape)
```

```
Training Features Shape: (261, 17)
Training Labels Shape: (261,)
Testing Features Shape: (87, 17)
Testing Labels Shape: (87,)
```

```
In [8]: # The baseline predictions are the historical averages
baseline_preds = test_features[:, feature_list.index('average')]
# Baseline errors, and display average baseline error
baseline_errors = abs(baseline_preds - test_labels)
print('Average baseline error: ', round(np.mean(baseline_errors), 2))
```

```
Average baseline error: 5.06
```

```
In [9]: # Import the model we are using
from sklearn.ensemble import RandomForestRegressor
# Instantiate model with 1000 decision trees
rf = RandomForestRegressor(n_estimators = 1000, random_state = 42)
# Train the model on training data
rf.fit(train_features, train_labels);
```

```
In [10]: # Use the forest's predict method on the test data
predictions = rf.predict(test_features)
# Calculate the absolute errors
errors = abs(predictions - test_labels)
# Print out the mean absolute error (mae)
print('Mean Absolute Error:', round(np.mean(errors), 2), 'degrees.')
```

```
Mean Absolute Error: 3.87 degrees.
```

```
In [11]: # Calculate mean absolute percentage error (MAPE)
mape = 100 * (errors / test_labels)
# Calculate and display accuracy
accuracy = 100 - np.mean(mape)
print('Accuracy:', round(accuracy, 2), '%.')
```

```
Accuracy: 93.93 %.
```

```
In [12]: # Import tools needed for visualization
from sklearn.tree import export_graphviz
import pydot
# Pull out one tree from the forest
tree = rf.estimators_[5]
# Import tools needed for visualization
from sklearn.tree import export_graphviz
import pydot
# Pull out one tree from the forest
tree = rf.estimators_[5]
# Export the image to a dot file
export_graphviz(tree, out_file = 'tree.dot', feature_names = feature_list, rounded = True, precision = 1)
# Use dot file to create a graph
(graph, ) = pydot.graph_from_dot_file('tree.dot')
# Write graph to a png file
graph.write_png('tree.png')
```


```
In [13]: # Limit depth of tree to 3 levels
rf_small = RandomForestRegressor(n_estimators=10, max_depth = 3)
rf_small.fit(train_features, train_labels)
# Extract the small tree
tree_small = rf_small.estimators_[5]
# Save the tree as a png image
export_graphviz(tree_small, out_file = 'small_tree.dot', feature_names = feature_list, rounded = True, precision = 1)
(graph, ) = pydot.graph_from_dot_file('small_tree.dot')
graph.write_png('small_tree.png');
```

```
In [14]: # Get numerical feature importances
importances = list(rf.feature_importances_)
# List of tuples with variable and importance
feature_importances = [(feature, round(importance, 2)) for feature, importance in zip(feature_list, importances)]
# Sort the feature importances by most important first
feature_importances = sorted(feature_importances, key = lambda x: x[1], reverse = True)
# Print out the feature and importances
[print('Variable: {:20} Importance: {}'.format(*pair)) for pair in feature_importances];
```

Variable: temp_1	Importance: 0.66
Variable: average	Importance: 0.15
Variable: forecast_noaa	Importance: 0.05
Variable: forecast_acc	Importance: 0.03
Variable: day	Importance: 0.02
Variable: temp_2	Importance: 0.02
Variable: forecast_under	Importance: 0.02
Variable: friend	Importance: 0.02
Variable: month	Importance: 0.01
Variable: year	Importance: 0.0
Variable: week_Fri	Importance: 0.0
Variable: week_Mon	Importance: 0.0
Variable: week_Sat	Importance: 0.0
Variable: week_Sun	Importance: 0.0
Variable: week_Thurs	Importance: 0.0
Variable: week_Tues	Importance: 0.0
Variable: week_Wed	Importance: 0.0

```
In [15]: # New random forest with only the two most important variables
rf_most_important = RandomForestRegressor(n_estimators= 1000, random_state=42)
# Extract the two most important features
important_indices = [feature_list.index('temp_1'), feature_list.index('average')]
train_important = train_features[:, important_indices]
test_important = test_features[:, important_indices]
# Train the random forest
rf_most_important.fit(train_important, train_labels)
# Make predictions and determine the error
predictions = rf_most_important.predict(test_important)
errors = abs(predictions - test_labels)
# Display the performance metrics
print('Mean Absolute Error:', round(np.mean(errors), 2), 'degrees.')
mape = np.mean(100 * (errors / test_labels))
accuracy = 100 - mape
print('Accuracy:', round(accuracy, 2), '%.')
```

Output



b) Implement Boosting ensemble method on a given dataset.

Algorithm

(b) Implement Boosting Ensemble on a given dataset.

Algorithm:

1. Initialize the dataset and assign equal weight to each of the data point.
2. Provide this as input to the model and identify the wrongly classified datapoints
3. increase the weight of the wrongly classified data points and decrease the weights of correctly classified data points. And then normalize the weights of all data points.
4. If (got required results)
 Goto step - 5
Else
 Goto step - 2
5. End.

Output

Confusion Matrix : $\begin{bmatrix} 116 & 35 \\ 26 & 54 \end{bmatrix}$

Accuracy score : 0.7359.

Code

```
In [1]: import pandas as pd
import numpy as np
import matplotlib.pyplot as plt
import seaborn as sns

%matplotlib inline
sns.set_style("whitegrid")
plt.style.use("fivethirtyeight")
```

```
In [2]: df = pd.read_csv("/content/diabetes.csv")
df.head()
```

```
Out[2]:   Pregnancies  Glucose  BloodPressure  SkinThickness  Insulin  BMI  DiabetesPedigreeFunction  Age  Outcome
0           6        148            72             35       0    33.6            0.627    50         1
1           1         85            66             29       0    26.6            0.351    31         0
2           8        183            64              0       0    23.3            0.672    32         1
3           1         89            66             23       94   28.1            0.167    21         0
4           0        137            40             35     168   43.1            2.288    33         1
```

```
In [3]: df.info()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 768 entries, 0 to 767
Data columns (total 9 columns):
 #   Column          Non-Null Count  Dtype  
 ---  --  
 0   Pregnancies      768 non-null   int64  
 1   Glucose          768 non-null   int64  
 2   BloodPressure    768 non-null   int64  
 3   SkinThickness    768 non-null   int64  
 4   Insulin          768 non-null   int64  
 5   BMI              768 non-null   float64 
 6   DiabetesPedigreeFunction 768 non-null   float64 
 7   Age              768 non-null   int64  
 8   Outcome          768 non-null   int64  
dtypes: float64(2), int64(7)
memory usage: 54.1 KB
```

```
In [4]: df.isnull().sum()
```

```
Out[4]: Pregnancies      0
Glucose          0
BloodPressure    0
SkinThickness    0
Insulin          0
BMI              0
DiabetesPedigreeFunction 0
Age              0
Outcome          0
dtype: int64
```

```
In [5]: pd.set_option('display.float_format', '{:.2f}'.format)
df.describe()
```

```
Out[5]:   Pregnancies  Glucose  BloodPressure  SkinThickness  Insulin  BMI  DiabetesPedigreeFunction  Age  Outcome
count    768.00     768.00     768.00      768.00    768.00    768.00            768.00    768.00    768.00
mean     3.85     120.89      69.11      20.54    79.80    31.99            0.47     33.24    0.35
std      3.37      31.97     19.36      15.95   115.24     7.88            0.33     11.76    0.48
min      0.00      0.00      0.00      0.00     0.00     0.00            0.08     21.00    0.00
25%     1.00     99.00      62.00      0.00     0.00    27.30            0.24     24.00    0.00
50%     3.00    117.00      72.00      23.00    30.50    32.00            0.37     29.00    0.00
75%     6.00    140.25      80.00      32.00   127.25    36.60            0.63     41.00    1.00
max     17.00    199.00     122.00     99.00   846.00   67.10            2.42     81.00    1.00
```

```
In [6]: categorical_val = []
continous_val = []
for column in df.columns:
#     print('====')
#     print(f'{column} : {df[column].unique()}')
    if len(df[column].unique()) <= 10:
        categorical_val.append(column)
    else:
        continous_val.append(column)

In [7]: df.columns

Out[7]: Index(['Pregnancies', 'Glucose', 'BloodPressure', 'SkinThickness', 'Insulin',
       'BMI', 'DiabetesPedigreeFunction', 'Age', 'Outcome'],
       dtype='object')

In [8]: # How many missing zeros are missing in each feature
feature_columns = [
    'Pregnancies', 'Glucose', 'BloodPressure', 'SkinThickness',
    'Insulin', 'BMI', 'DiabetesPedigreeFunction', 'Age'
]

for column in feature_columns:
    print("====")
    print(f'{column} ==> Missing zeros : {len(df.loc[df[column] == 0])}')

=====
Pregnancies ==> Missing zeros : 111
=====
Glucose ==> Missing zeros : 5
=====
BloodPressure ==> Missing zeros : 35
=====
SkinThickness ==> Missing zeros : 227
=====
Insulin ==> Missing zeros : 374
=====
BMI ==> Missing zeros : 11
=====
DiabetesPedigreeFunction ==> Missing zeros : 0
=====
Age ==> Missing zeros : 0

In [9]: from sklearn.impute import SimpleImputer

fill_values = SimpleImputer(missing_values=0, strategy="mean", copy=False)
df[feature_columns] = fill_values.fit_transform(df[feature_columns])

for column in feature_columns:
    print("====")
    print(f'{column} ==> Missing zeros : {len(df.loc[df[column] == 0])}')

=====
Pregnancies ==> Missing zeros : 0
=====
Glucose ==> Missing zeros : 0
=====
BloodPressure ==> Missing zeros : 0
=====
SkinThickness ==> Missing zeros : 0
=====
Insulin ==> Missing zeros : 0
=====
BMI ==> Missing zeros : 0
=====
DiabetesPedigreeFunction ==> Missing zeros : 0
=====
Age ==> Missing zeros : 0
```

```
In [10]: from sklearn.model_selection import train_test_split

X = df[feature_columns]
y = df.Outcome

X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.3, random_state=42)
```

```
In [11]: from sklearn.metrics import confusion_matrix, accuracy_score, classification_report

def evaluate(model, X_train, X_test, y_train, y_test):
    y_test_pred = model.predict(X_test)
    y_train_pred = model.predict(X_train)

    print("TRAINING RESULTS: \n====")
    clf_report = pd.DataFrame(classification_report(y_train, y_train_pred, output_dict=True))
    print(f"CONFUSION MATRIX:\n{confusion_matrix(y_train, y_train_pred)}")
    print(f"ACCURACY SCORE:\n{accuracy_score(y_train, y_train_pred):.4f}")
    print(f"CLASSIFICATION REPORT:\n{clf_report}")

    print("TESTING RESULTS: \n====")
    clf_report = pd.DataFrame(classification_report(y_test, y_test_pred, output_dict=True))
    print(f"CONFUSION MATRIX:\n{confusion_matrix(y_test, y_test_pred)}")
    print(f"ACCURACY SCORE:\n{accuracy_score(y_test, y_test_pred):.4f}")
    print(f"CLASSIFICATION REPORT:\n{clf_report}")
```

```
In [12]: from sklearn.ensemble import AdaBoostClassifier

ada_boost_clf = AdaBoostClassifier(n_estimators=30)
ada_boost_clf.fit(X_train, y_train)
evaluate(ada_boost_clf, X_train, X_test, y_train, y_test)
```

Output-AdaBoost

```
TRAINING RESULTS:
=====
CONFUSION MATRIX:
[[310  39]
 [ 51 137]]
ACCURACY SCORE:
0.8324
CLASSIFICATION REPORT:
          0      1  accuracy  macro avg  weighted avg
precision  0.86   0.78      0.83      0.82      0.83
recall    0.89   0.73      0.83      0.81      0.83
f1-score   0.87   0.75      0.83      0.81      0.83
support   349.00 188.00      537.00      537.00
TESTING RESULTS:
=====
CONFUSION MATRIX:
[[123  28]
 [ 27  53]]
ACCURACY SCORE:
0.7619
CLASSIFICATION REPORT:
          0      1  accuracy  macro avg  weighted avg
precision  0.82   0.65      0.76      0.74      0.76
recall    0.81   0.66      0.76      0.74      0.76
f1-score   0.82   0.66      0.76      0.74      0.76
support   151.00  80.00      231.00      231.00
```

Output- GradientBoost


```
TRAINING RESULTS:  
=====  
CONFUSION MATRIX:  
[[342  7]  
 [ 19 169]]  
ACCURACY SCORE:  
0.9516  
CLASSIFICATION REPORT:  
          0      1  accuracy  macro avg  weighted avg  
precision  0.95  0.96      0.95      0.95      0.95  
recall    0.98  0.90      0.95      0.94      0.95  
f1-score   0.96  0.93      0.95      0.95      0.95  
support   349.00 188.00      537.00      537.00  
TESTING RESULTS:  
=====  
CONFUSION MATRIX:  
[[116  35]  
 [ 26  54]]  
ACCURACY SCORE:  
0.7359  
CLASSIFICATION REPORT:  
          0      1  accuracy  macro avg  weighted avg  
precision  0.82  0.61      0.74      0.71      0.74  
recall    0.77  0.68      0.74      0.72      0.74  
f1-score   0.79  0.64      0.74      0.72      0.74  
support   151.00 80.00      231.00      231.00
```

PROGRAM 10

Date: 24-05-2024

Build k-Means algorithm to cluster a set of data stored in a .CSV file.

Algorithm



Code

Importing and initializing the data points


```
In [5]: import matplotlib.pyplot as plt  
import numpy as np  
from sklearn.cluster import KMeans  
  
In [6]: from sklearn.datasets import make_blobs  
X, y_true = make_blobs(n_samples=550, centers=4, cluster_std=0.6, random_state=0)  
  
In [7]: import plotly.express as px  
fig = px.scatter(x=X[:, 0], y=X[:, 1], width=800, height=500)  
fig.show()
```

Elbow Method to find optimal K

```
In [8]: cost = []  
for i in range(1, 11):  
    km = KMeans(n_clusters=i, max_iter=500)  
    km.fit(X)  
  
    cost.append(km.inertia_)  
  
# plot the cost against K values  
fig = px.line(x=range(1, 11), y=cost, width=600, height=400)  
fig.show()  
# the point of the elbow is the  
# most optimal value for choosing k
```

Defining Model and fitting the same

```
In [9]: kmeans = KMeans(n_clusters=4)  
kmeans.fit(X)  
y_kmeans = kmeans.predict(X)  
  
In [10]: fig = px.scatter(x=X[:, 0], y=X[:, 1], color=y_kmeans, width=700, height=400)  
trace = px.scatter(x=X[:, 0], y=X[:, 1], width=700, height=400)  
fig.show()  
  
In [11]: plt.scatter(X[:, 0], X[:, 1], c=y_kmeans, s=20)  
centers = kmeans.cluster_centers_  
plt.scatter(centers[:, 0], centers[:, 1], c='blue', s=100, alpha=0.9);  
plt.show()
```



Iris Dataset

```
In [12]: import pandas as pd
import seaborn as sns
import matplotlib.pyplot as plt
from sklearn import datasets

In [24]: iris = datasets.load_iris()
df = pd.DataFrame(iris.data)
df['class']=iris.target
df.columns=['sepal_len', 'sepal_wid', 'petal_len', 'petal_wid', 'class']
df.info()

<class 'pandas.core.frame.DataFrame'>
RangeIndex: 150 entries, 0 to 149
Data columns (total 5 columns):
 #   Column      Non-Null Count  Dtype  
--- 
 0   sepal_len    150 non-null    float64
 1   sepal_wid    150 non-null    float64
 2   petal_len    150 non-null    float64
 3   petal_wid    150 non-null    float64
 4   class        150 non-null    int64  
dtypes: float64(4), int64(1)
memory usage: 6.0 KB

In [52]: px.histogram(df, x ='class', color='class')

In [56]: from sklearn.preprocessing import StandardScaler
scaler = StandardScaler()
X = df.iloc[:,0:4].values

In [63]: scaled_x = scaler.fit_transform(X)

In [74]: model = KMeans(n_clusters=3,init='k-means++',random_state=0)
labels = model.fit_predict(scaled_x)

In [81]: import plotly.graph_objects as go
fig = go.Figure()


# Add trace
fig.add_trace(go.Histogram(x=labels,name="Predicted Labels"))
fig.add_trace(go.Histogram(x=df['class'],name="True Labels"))

# Overlay both histograms
fig.update_layout(barmode='overlay')
# Reduce opacity to see both histograms
fig.update_traces(opacity=0.75)
fig.show()

In [89]: labels =[]
for i in range(2, 5):
    model = KMeans(n_clusters = i, max_iter = 500)
    model.fit(scaled_x)
    labels.append(model.fit_predict(scaled_x))

In [97]: from plotly.subplots import make_subplots
import plotly.graph_objects as go
fig = make_subplots(rows=2, cols=2)
for i in range(0, 3):
    fig.add_trace(go.Histogram(x=labels[i],name="{} Clusters".format(i+2)),
                 row=(i//2 + 1), col=(i%2 + 1))
fig.add_trace(go.Histogram(x=df['class'],name="True Classification"),
             row=2, col=2)
fig.update_layout(height=700, width=1000, title_text="Side By Side Subplots")
fig.show()
```


Output



PROGRAM 11

Date: 24-05-2024

Implement Dimensionality reduction using Principle Component Analysis (PCA) method.



Code

```
Open in Colab

In [ ]: from google.colab import drive
drive.mount('/content/drive')

Mounted at /content/drive

In [ ]: import seaborn as sns
import numpy as np
import pandas as pd
import matplotlib.pyplot as plt
import plotly.express as px
import plotly.graph_objects as go
from plotly.subplots import make_subplots

In [ ]: df = pd.read_csv('/content/drive/MyDrive/breast-cancer.csv')
df.head()

Out[ ]:
      id diagnosis radius_mean texture_mean perimeter_mean area_mean smoothness_mean compactness_mean concavity_mean
0   842302        M       17.99      10.38       122.80     1001.0      0.11840      0.27760      0.3001
1   842517        M       20.57      17.77       132.90     1326.0      0.08474      0.07864      0.0869
2   84300903       M       19.69      21.25       130.00     1203.0      0.10960      0.15990      0.1974
3   84348301       M       11.42      20.38       77.58      386.1      0.14250      0.28390      0.2414
4   84358402       M       20.29      14.34       135.10     1297.0      0.10030      0.13280      0.1980

5 rows × 32 columns

In [ ]: df.drop('id', axis=1, inplace=True) #drop redundant columns

In [ ]: df['diagnosis'] = (df['diagnosis'] == 'M').astype(int) #encode the Label into 1/0

In [ ]: corr = df.corr()

In [ ]:
# Get the absolute value of the correlation
cor_target = abs(corr["diagnosis"])

# Select highly correlated features (threshold = 0.2)
relevant_features = cor_target[cor_target>0.2]

# Collect the names of the features
names = [index for index, value in relevant_features.items()]

# Drop the target variable from the results
names.remove('diagnosis')

# Display the results
print(names)

['radius_mean', 'texture_mean', 'perimeter_mean', 'area_mean', 'smoothness_mean', 'compactness_mean', 'concavity_mean', 'concave points_mean', 'symmetry_mean', 'radius_se', 'perimeter_se', 'area_se', 'compactness_se', 'concavity_se', 'concave points_se', 'radius_worst', 'texture_worst', 'perimeter_worst', 'area_worst', 'smoothness_worst', 'compactness_worst', 'concavity_worst', 'concave points_worst', 'symmetry_worst', 'fractal_dimension_worst']

In [ ]: X = df[names].values
```

```

In [ ]:
class PCA:
    """
    Principal Component Analysis (PCA) class for dimensionality reduction.
    """

    def __init__(self, n_components):
        """
        Constructor method that initializes the PCA object with the number of components to retain.

        Args:
        - n_components (int): Number of principal components to retain.
        """
        self.n_components = n_components

    def fit(self, X):
        """
        Fits the PCA model to the input data and computes the principal components.

        Args:
        - X (numpy.ndarray): Input data matrix with shape (n_samples, n_features).
        """
        # Compute the mean of the input data along each feature dimension.
        mean = np.mean(X, axis=0)

        # Subtract the mean from the input data to center it around zero.
        X = X - mean

        # Compute the covariance matrix of the centered input data.
        cov = np.cov(X.T)

        # Compute the eigenvectors and eigenvalues of the covariance matrix.
        eigenvalues, eigenvectors = np.linalg.eigh(cov)
        # Reverse the order of the eigenvalues and eigenvectors.
        eigenvalues = eigenvalues[::-1]
        eigenvectors = eigenvectors[:, ::-1]

        # Keep only the first n_components eigenvectors as the principal components.
        self.components = eigenvectors[:, :self.n_components]

        # Compute the explained variance ratio for each principal component.
        # Compute the total variance of the input data
        total_variance = np.sum(np.var(X, axis=0))

        # Compute the variance explained by each principal component
        self.explained_variances = eigenvalues[:self.n_components]

        # Compute the explained variance ratio for each principal component
        self.explained_variance_ratio_ = self.explained_variances / total_variance

    def transform(self, X):
        """
        Transforms the input data by projecting it onto the principal components.

        Args:
        - X (numpy.ndarray): Input data matrix with shape (n_samples, n_features).

        Returns:
        - transformed_data (numpy.ndarray): Transformed data matrix with shape (n_samples, n_components).
        """
        # Center the input data around zero using the mean computed during the fit step.
        X = X - np.mean(X, axis=0)

        # Project the centered input data onto the principal components.
        transformed_data = np.dot(X, self.components)

        return transformed_data

    def fit_transform(self, X):
        """
        Fits the PCA model to the input data and computes the principal components then
        transforms the input data by projecting it onto the principal components.

        Args:
        - X (numpy.ndarray): Input data matrix with shape (n_samples, n_features).
        """
        self.fit(X)
        transformed_data = self.transform(X)
        return transformed_data

```

```
In [ ]: pca = PCA(2)

In [ ]: pca.fit(X)

In [ ]: pca.explained_variance_ratio_

Out[ ]: array([0.98377428, 0.01620498])

In [ ]: X_transformed = pca.transform(X)

In [ ]: X_transformed[:,1].shape

Out[ ]: (569,)

In [ ]: fig = px.scatter(x=X_transformed[:,0], y=X_transformed[:,1])
fig.update_layout(
    title="PCA transformed data for breast cancer dataset",
    xaxis_title="PC1",
    yaxis_title="PC2"
)
fig.show()
```

Output

